

# Collaborative Filtering with Matrix Factorization: Book Recommendation System

Hardcover Book Friend Finder

January 2, 2026

## Abstract

We present a hybrid book recommendation system using collaborative filtering with matrix factorization. The model decomposes a sparse user-book interaction matrix into two lower-dimensional matrices representing latent features. We employ masking to handle extreme sparsity (95.75%) and combine collaborative filtering with popularity-based recommendations to achieve 8.83% precision@10, outperforming pure collaborative filtering by 308%.

## 1 Matrix Factorization

### 1.1 Problem Setup

Given a user-book interaction matrix  $Y \in \mathbb{R}^{m \times n}$  where:

- $m =$  number of books (2,547)
- $n =$  number of users (246)
- $Y_{ij}$  represents user  $j$ 's interaction with book  $i$

The matrix is highly sparse: only 26,598 interactions out of 626,562 possible (95.75% sparse).

### 1.2 Data Filtering

From the raw dataset of 1,000 users and 45,203 books, we applied filtering to improve model quality:

- **User filter:** Keep only users with  $\geq 20$  book interactions (246 users retained, 24.6%)
- **Book filter:** Keep only books with  $\geq 5$  users (2,547 books retained, 5.6%)

This reduces noise from inactive users and obscure books while maintaining sufficient data for collaborative filtering. The filtered dataset provides better signal for learning user preferences.


### 1.3 Matrix Decomposition

We factor  $Y$  into two rectangular matrices:

$$Y \approx \sigma(XW^T + b) \tag{1}$$

where:

- $X \in \mathbb{R}^{m \times k}$  is the **book feature matrix** ( $2,547 \times 20$ )
- $W \in \mathbb{R}^{n \times k}$  is the **user feature matrix** ( $246 \times 20$ )
- $b \in \mathbb{R}^{1 \times n}$  is the user bias vector
- $k = 20$  latent features
- $\sigma$  is the sigmoid activation function



## 1.4 Latent Features

The  $k = 20$  latent features are **learned representations** that capture:

- **Book features ( $X$ ):** Genre preferences (fantasy, sci-fi, literary fiction), reading level, book popularity, themes
- **User features ( $W$ ):** User's affinity for each latent dimension, reading preferences, taste profile

These features are **not predefined** – they emerge automatically during training through gradient descent. Each row of  $W$  becomes a 20-dimensional embedding representing a user's reading preferences in latent space.

## 1.5 Prediction

For user  $j$  and book  $i$ :

$$\hat{y}_{ij} = \sigma \left( \sum_{f=1}^k x_{if} w_{jf} + b_j \right) \quad (2)$$

This predicts the probability that user  $j$  will like book  $i$ .

## 2 Handling Sparsity with Masking

### 2.1 The Masking Matrix

Define indicator matrix  $R \in \{0, 1\}^{m \times n}$ :

$$R_{ij} = \begin{cases} 1 & \text{if user } j \text{ has interacted with book } i \\ 0 & \text{otherwise} \end{cases} \quad (3)$$

## 2.2 Rating Scale with Implicit Feedback

$$Y_{ij} = \begin{cases} 1.0 & \text{read with rating } \geq 3 \\ 0.7 & \text{currently reading} \\ 0.3 & \text{want to read} \\ 0.0 & \text{read with rating } < 3 \text{ or DNF} \\ 0.5 & \text{unread (placeholder, ignored)} \end{cases} \quad (4)$$

## 2.3 Masked Loss Function

The cost function only considers **known interactions**:

$$J(X, W, b) = \sum_{i=1}^m \sum_{j=1}^n R_{ij} (\sigma(x_i^T w_j + b_j) - Y_{ij})^2 + \frac{\lambda}{2} (\|X\|_F^2 + \|W\|_F^2) \quad (5)$$

where:

- $R_{ij}$  masks out unread books (where  $Y_{ij} = 0.5$ )
- $\lambda = 1.0$  is the regularization parameter
- $\|\cdot\|_F$  is the Frobenius norm

**Key insight:** Without masking, the 95.75% of unread books would dominate the loss and the model would simply learn the global average.

## 3 Training

### 3.1 Optimization

We use Adam optimizer with learning rate 0.1 for 300 iterations:

$$\nabla_X J = R \odot [(\sigma(XW^T + b) - Y) \odot \sigma'(XW^T + b)] W + \lambda X \quad (6)$$

$$\nabla_W J = R^T \odot [(\sigma(XW^T + b) - Y) \odot \sigma'(XW^T + b)]^T X + \lambda W \quad (7)$$

$$\nabla_b J = \sum_{i=1}^m R_{ij} (\sigma(x_i^T w_j + b_j) - Y_{ij}) \sigma'(x_i^T w_j + b_j) \quad (8)$$

where  $\odot$  denotes element-wise multiplication and  $\sigma'(z) = \sigma(z)(1 - \sigma(z))$ .

### 3.2 Training Time

Total compilation: 4.72 seconds

- Data loading: 1.40s
- Matrix building: 0.02s
- Model training: 3.26s
- User clustering: 0.04s

## 4 Hybrid Model

Pure collaborative filtering achieved only 2.45% precision@10. We combine with popularity:

$$\text{score}_{ij} = 0.5 \cdot \text{popularity}_i + 0.5 \cdot \sigma(x_i^T w_j + b_j) \quad (9)$$

where  $\text{popularity}_i = \frac{\sum_j R_{ij}}{\max_i \sum_j R_{ij}}$  is the normalized book popularity.

## 5 Results

Method	Precision@10
Pure collaborative filtering	2.16%
Improved collaborative (filtered)	2.45%
Popularity baseline	7.56%
<b>Hybrid (50/50)</b>	<b>8.83%</b>

Table 1: Recommendation accuracy on 80/20 train/test split. Hybrid model achieves 308% improvement over pure collaborative filtering and 17% improvement over popularity baseline.

### 5.1 Why Hybrid Works

- **Popularity component (50%)**: Ensures quality by recommending well-regarded books
- **Collaborative component (50%)**: Adds personalization based on user’s latent preferences
- **Implicit feedback**: Increased training data by 75% (`want_to_read` + `currently_reading`)

## 6 Comparison: Netflix vs. Hardcover

Aspect	Netflix	Hardcover
Rating scale	{−1, 1, 0} (dislike, like, unranked)	{0, 0.3, 0.7, 1, 0.5} (implicit feedback)
Unranked handling	0 = unranked, gets predicted	0.5 = placeholder, <b>masked out</b>
Activation	Linear (outputs $\in \{-1, 1\}$ )	Sigmoid (outputs $\in [0, 1]$ )
Loss computation	All entries	<b>Only known</b> (R=1)
Sparsity	~95-98%	95.75%
Data density	Dense enough to predict 0s	Too sparse, must mask

Table 2: Key differences between Netflix-style and Hardcover collaborative filtering approaches.

## 7 Clustering for Friend Matching

### 7.1 Determining Optimal Cluster Count

We apply K-means clustering on L2-normalized user embeddings  $W$  to group users with similar reading preferences. To determine the optimal number of clusters, we tested  $K \in [3, 15]$  and

evaluated using:

- **Silhouette score:** Measures cluster separation (higher is better)
- **Calinski-Harabasz score:** Ratio of between-cluster to within-cluster variance
- **Elbow method:** Identifies diminishing returns in inertia reduction

The silhouette score was maximized at  $K = 13$  clusters (score: 0.084), indicating this provides the best balance between cluster cohesion and separation.

## 7.2 Friend Matching within Clusters

Users within the same cluster are ranked by cosine similarity:

$$\text{similarity}(u_i, u_j) = \frac{w_i \cdot w_j}{\|w_i\| \|w_j\|} \quad (10)$$

where  $w_i, w_j \in \mathbb{R}^{20}$  are the 20-dimensional user embeddings (rows of  $W$ ). Note: we use  $k = 20$  latent **features** for matrix factorization, but group users into  $K = 13$  **clusters** for friend matching.

The 13 clusters represent distinct reading groups such as "Classic YA & Fantasy Fans," "Sci-Fi Enthusiasts," and "Literary Fiction Lovers," with cluster sizes ranging from 4 to 52 users.

## 8 Conclusion

Matrix factorization with masking effectively handles sparse user-book interactions. The key innovations are:

1. **Masking** unread books to avoid learning from noise
2. **Implicit feedback** to increase training signals by 75%
3. **Hybrid ensemble** balancing popularity and personalization
4. **Low-dimensional embeddings** ( $k = 20$ ) capturing user preferences

The system achieves 8.83% precision@10, providing personalized book recommendations for 246 active readers.